

COMP9334: Capacity Planning of Computer Systems and Networks

Week 10: Optimisation (1)

A/Prof Chun Tung Chou
CSE, UNSW

Three Weeks of Optimisation


- The lectures for these three weeks will focus on optimization methods for network related design and applications
 - Week 10: Linear programming, integer programming
 - Weeks 11-12: More applications of integer programming
- You will learn:
 - How to formulate optimization problems
 - Tools to solve optimization problems
- An introduction only, because optimization is a big topic
 - Emphasis is on applying optimization methods rather than the theory behind

Motivation (1)

- A modern approach to managing computer networks is based on the concept of *software-defined networking*
- Two types of nodes:
 1. Simple packet switches
 2. Controllers
- A controller can control a number of simple packet switches but they must be placed in a strategic location in the network
- If the delay between the controller and a packet switch is too long, then it can degrade the network performance

Motivation (2)

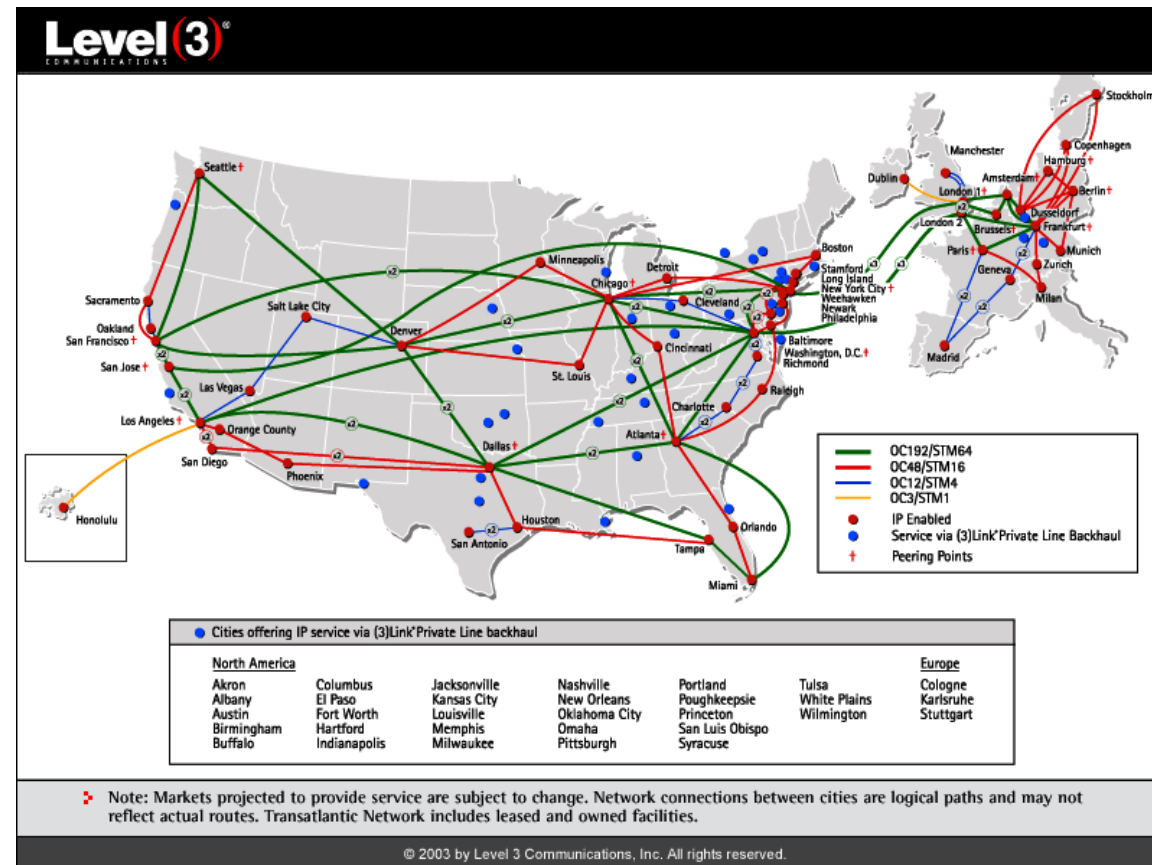
- Consider the following network where there is a packet switch at each node and the delay on each link is 1 time unit.
- Question: Assuming you want to place one controller in the network, where will you place the controller?



- Question: What if you want to place two controllers?

Motivation (3)

- How about solving the same problem for a large heterogeneous network?



- Optimisation provides a systematic method to make decisions

Elements of an optimisation problem

- You want to maximise your WAM and still have a life

Maximise $WAM(x_1, x_2, x_3, \dots)$

x_1 hours/week on COMP9334

x_2 hours/week on COMPxxxx

x_3 hours/week on socialising

$x_1 \geq 10$

$x_2 \leq \text{maxSocialHours}$

$x_1 + x_2 + x_3 + \dots \leq \text{totalAwakeHours}$

- Elements of an optimisation problem
 - Minimise or maximise an objective function
 - Decision variables: x_1, x_2, \dots etc.
 - Constraints

What is optimization?

- In mathematics, also known as **mathematical programming**
 - The term **programming** refers to planning of activities to obtain an optimal result, not computer programming
 - The amount or level of each activity can be represented as a variable whose value is to be determined
- **Optimization** means solving problems in which we seek to minimize or maximize the value of an **objective function** of many **decision variables**, subject to **constraints** on the decision variables

Reference books

- Winston, “Operations Research”, 4th edition
 - Examples from this book tend to come from manufacturing, business, finance, etc
 - The abstraction power of mathematics means many optimization problems have similar mathematical formulation
 - Very often an optimization problem in networking may have a similar cousin in other application areas, and their mathematical formulation are identical
- Ahuja, Magnanti and Orlin, “Network Flows”
- Fourer, Gay and Kernighan, “AMPL: A Modeling Language for Mathematical Programming”, 2nd edition


Software

- Modeling language: AMPL and Solver CPLEX
 - High-level programming language for describing optimization problems
 - Syntax similar to mathematical formulation of optimization problems
 - Demo version of the software is available for download from:
<http://www.ampl.com>. Click *Try AMPL*, then *Download a Free Demo*
- Note: Demo version of AMPL/CPLEX is full-featured but limited to 500 variables and 500 objectives plus constraints


Motivating example 1: Cloud/Grid computing

- Service providers sell computing power as an utility
 - Computing power measured in CPU cycles
- Target customers
 - Financial company, pharmaceutical company, etc.
- Quality of Service in Cloud computing
 - Different service providers might offer the service at different levels for different costs
 - Optimization problem: How to select service providers (allocate resources) to achieve the best level of service without exceeding budget


Cloud computing resource allocation



Resource 1
Speed: 1,000 million
cycles/sec
Cost: 0.1 dollars/sec




Resource 2
Speed: 2,000 million
cycles/sec
Cost: 0.25 dollars/sec




Resource 3
Speed: 3,000 million
cycles/sec
Cost: 0.6 dollars/sec

- A computation job:
 - Requires 10^7 million cycles
 - Must be completed in at most 4,800 sec
 - Cost must not exceed 1,500 dollars
- Exercises: For the time being, let us ignore the constraint on the completion time and cost.
 - If you use Resource 1 only, what is the completion time and cost?
 - Repeat for Resources 2 and 3.


Cloud computing resource allocation (cont.)



Resource 1
Speed: 1,000 million
cycles/sec
Cost: 0.1 dollars/sec




Resource 2
Speed: 2,000 million
cycles/sec
Cost: 0.25 dollars/sec




Resource 3
Speed: 3,000 million
cycles/sec
Cost: 0.6 dollars/sec

- A computation job:
 - Requires 10^7 million cycles
 - Must be completed in at most 4,800 sec
 - Cost must not exceed 1,500 dollars
- Completion time and cost for each resource:
 - Resource 1: ?
 - Resource 2: ?
 - Resource 3: ?


Cloud computing resource allocation (cont.)



Resource 1
Speed: 1,000 million
cycles/sec
Cost: 0.1 dollars/sec



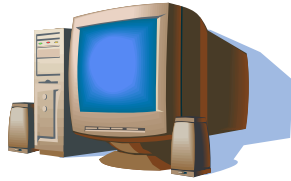
Resource 2
Speed: 2,000 million
cycles/sec
Cost: 0.25 dollars/sec



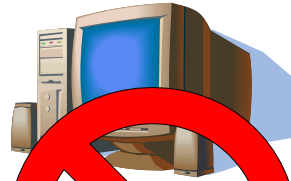
Resource 3
Speed: 3,000 million
cycles/sec
Cost: 0.6 dollars/sec

- Assume the computation job can be arbitrarily split into up to three parallel tasks
- Question: How should the job be split, so that completion time T is minimized subject to two constraints:
 - Completion time constraint: $T \leq 4,800$ sec
 - Cost constraint: $C \leq 1,500$ dollars

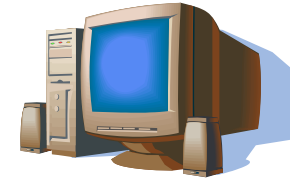
Trial and error: Solution 1



Resource 1
Speed: 1,000 million
cycles/sec
Cost: 0.1 dollars/sec



Resource 2
Speed: 2,000 million
cycles/sec
Cost: 0.25 dollars/sec




Resource 3
Speed: 3,000 million
cycles/sec
Cost: 0.6 dollars/sec

- 48% to Resource 1, 52% to Resource 3
 - Resource 1: Completion time = 4,800 sec, cost = 480 dollars
 - Resource 3: Completion time = 1,733 sec, cost = 1,040 dollars
- Job completion time = 4,800 sec (remember jobs run in parallel)
- cost = 1,520 dollars, **Infeasible** solution


Terminology

- A solution is **feasible** if all the constraints are satisfied
- A solution is **infeasible** if not all the constraints are satisfied


Trial and error: Solution 2



Resource 1
Speed: 1,000 million
cycles/sec
Cost: 0.1 dollars/sec




Resource 2
Speed: 2,000 million
cycles/sec
Cost: 0.25 dollars/sec




Resource 3
Speed: 3,000 million
cycles/sec
Cost: 0.6 dollars/sec

- 70% to Resource 2, 30% to Resource 3
 - Resource 2: Completion time = 3,500 sec, cost = 875 dollars
 - Resource 3: Completion time = 1,000 sec, cost = 600 dollars
- Job completion time = 3,500 sec, cost = 1,475 dollars
- Feasible solution


Trial and error: Solution 3



Resource 1
Speed: 1,000 million
cycles/sec
Cost: 0.1 dollars/sec



Resource 2
Speed: 2,000 million
cycles/sec
Cost: 0.25 dollars/sec



Resource 3
Speed: 3,000 million
cycles/sec
Cost: 0.6 dollars/sec

- 30% to Resource 1, 30% to Resource 2, 40% to Resource 3
- Exercise: Check if this is feasible

Optimizing resource allocation

■ Given:

- Job requirement = 10^7 million cycles
- Completion time $\leq 4,800$ sec
- Budget $\leq 1,500$ dollars

■ Let:

- x_1 = fraction of the job to Resource 1
- x_2 = fraction of the job to Resource 2
- x_3 = fraction of the job to Resource 3

■ Find x_1 , x_2 and x_3 such that

- All requirements are met
- Completion time is minimized

Formulating optimization problem

■ Completion time:

- Resource 1 = $\frac{10^7 \times x_1}{1000} = 10000 \times x_1$

- Resource 2 = $\frac{10^7 \times x_2}{2000} = 5000 \times x_2$

- Resource 3 = $\frac{10^7 \times x_3}{3000} = \frac{10000}{3} \times x_3$

- Job completion time $T = \max(10000 \times x_1, 5000 \times x_2, \frac{10000}{3} \times x_3)$

■ Cost:

- Resource 1 = $0.1 \times 10000 \times x_1 = 1000 \times x_1$

- Resource 2 = $0.25 \times 5000 \times x_2 = 1250 \times x_2$

- Resource 3 = $0.6 \times \frac{10000}{3} \times x_3 = 2000 \times x_3$

- Cost $C = 1000 \times x_1 + 1250 \times x_2 + 2000 \times x_3$

Formulating optimization problem (cont.)

- Mathematically, the optimization problem can be formulated as

$$\min T$$

subject to

$$T \geq 10000 \times x_1$$

$$T \geq 5000 \times x_2$$

$$T \geq \frac{10000}{3} \times x_3$$

$$T \leq 4800$$

$$1000 \times x_1 + 1250 \times x_2 + 2000 \times x_3 \leq 1500$$

$$x_1 + x_2 + x_3 = 1$$

$$x_1, x_2, x_3 \geq 0$$

Components of an optimization problem

- Given parameters
- Decision variables
 - In this example, they are x_1 , x_2 , x_3 and T
- Objective function
 - Can be minimization or maximization
 - Can be single objective or multi-objective
- Constraints

Exercise

- Consider the following optimization problem where x is the decision variable:

$$\min_x 2x - 1$$

subject to

$$x \leq 20$$

$$x \geq 8$$


- What are the feasible solutions?
- What is the optimal solution?

Linear programming

- Linear programming (LP) is a tool for solving a type of optimization problems where
 - Decision variables are real numbers
 - Objective function is linear in the decision variables
 - All the constraints are linear in the decision variables
- LP problems with 2 variables have a nice graphical solution, e.g.


$$\begin{array}{ll}\min & x_1 + x_2 \\ \text{subject to} & 2x_1 + x_2 \geq 4 \\ & x_1 + 3x_2 \geq 6 \\ & x_1, x_2 \geq 0\end{array}$$

Feasible solutions




- Any (x_1, x_2) that satisfies all the constraints is a **feasible** solution. Otherwise, it is an **infeasible** solution
- For LP problems with 2 variables, the feasible region is a polygon
- In general, it is a polyhedron

Objective function




- Since the aim is to minimize $x_1 + x_2$, we look for contour of $x_1 + x_2$

Optimal solution



- Optimal solution of this LP problem is $x_1 = 1.2$, $x_2 = 1.6$, with objective function value = 2.8

Extreme point



- Regardless of the objective function, the optimal solution of an LP problem (if it exists) must be at one of the extreme points or vertices of the polyhedron formed by the constraints

Special cases: Multiple optimal solutions

- What is the optimal solution to the following LP problem?

$$\begin{array}{ll}\min & x_1 + x_2 \\ \text{subject to} & x_1 + x_2 \geq 4 \\ & x_1 + 3x_2 \geq 6 \\ & x_1, x_2 \geq 0\end{array}$$

Special cases: Infeasible LP

- What is the optimal solution to the following LP problem?

$$\begin{array}{ll}\min & x_1 + x_2 \\ \text{subject to} & x_1 + 3x_2 \leq 4 \\ & x_1 + 3x_2 \geq 6 \\ & x_1, x_2 \geq 0\end{array}$$


Special cases: Unbounded LP

- What is the optimal solution to the following LP problem?

$$\begin{array}{ll}\max & x_1 + x_2 \\ \text{subject to} & 2x_1 + x_2 \geq 4 \\ & x_1 + 3x_2 \geq 6 \\ & x_1, x_2 \geq 0\end{array}$$


Algorithms for LP

- Simplex method (Dantzig 1947)
 - Found in many LP software
 - Principle: Move from an extreme point to another
 - Worst-case complexity: Exponential
 - Average performance on practical LP problems: Very good



Algorithms for LP (cont.)

- Interior-point method (Karmarkar 1984)
 - Very competitive compared with the simplex method
 - Principle: Move from an interior point to another
 - Worst-case complexity: Polynomial
 - The algorithm for large LP problems



LP solvers

- Many commercial and free software are available for solving LP problems
- Commercial software
 - Capable of solving large LP problems, e.g. millions of variables
 - A 50,000-variable LP problem takes about 5 seconds on a standard linux PC
 - You can try out many commercial solvers at the NEOS web site
 - `https://neos-server.org/neos/`
- Free software / demo version
 - `http://www.ampl.com`
 - `http://ampl.com/try-ampl/download-a-demo-version/`

LP solvers (cont.)

- LP solvers require the user to write the problem in fixed format
- Can be embedded in C, C++ or Java, e.g.

```
model.add(IloMinimize(env, -9*x[0]+x[1]+4*x[2]));  
model.add(-x[0]+x[2] == -3);  
model.add( x[0]-x[1] <= 1);
```

- Can be used with some modeling languages
 - AMPL
 - MPS
 - GAMS

AMPL/CPLEX for solving example 1

- In AMPL, the grid computing problem formulated earlier becomes

```
var T;  
var x_1 >= 0;  
var x_2 >= 0;  
var x_3 >= 0;  
minimize time: T;  
subject to T_1: T >= 10000*x_1;  
subject to T_2: T >= 5000*x_2;  
subject to T_3: T >= 10000/3*x_3;  
subject to T_max: T <= 4800;  
subject to C_max: 1000*x_1+1250*x_2+2000*x_3 <= 1500;  
subject to x_sum: x_1+x_2+x_3 = 1;
```

- This is saved in the file `grid_lp.mod`

AMPL/CPLEX for solving example 1 (cont.)

- The problem can be solved by CPLEX with the batch file `grid_lp_batch`

```
model grid_lp.mod;  
option solver cplex;  
solve;  
display x_1;  
display x_2;  
display x_3;  
display T;
```

AMPL/CPLEX for solving example 1 (Command line version)

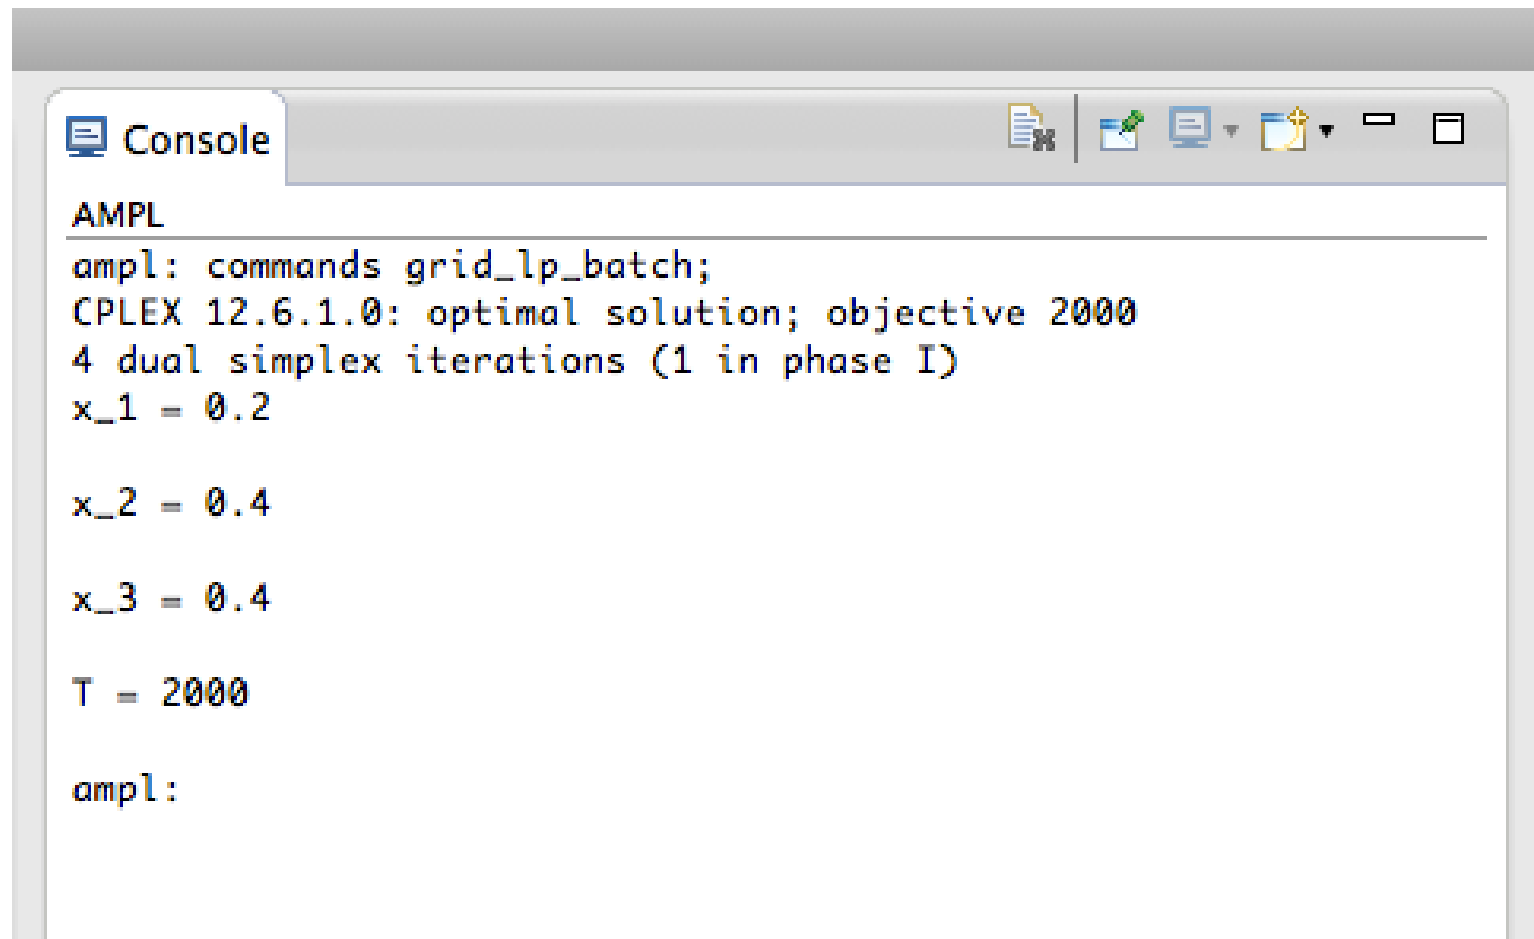
- At the ampl command prompt, type `commands grid_lp_batch;`, it returns

```
commands grid_lp_batch;  
CPLEX 12.6.0.0: optimal solution; objective 2000  
4 dual simplex iterations (1 in phase I)  
x_1 = 0.2  
  
x_2 = 0.4  
  
x_3 = 0.4  
  
T = 2000  
  
1000*x_1 + 1250*x_2 + 2000*x_3 = 1500
```

- Note: All these files can be downloaded from the course web site

AMPL/CPLEX for solving example 1 (IDE version)

- At the AMPL prompt, type `commands grid_lp_batch;`
- Need to `reset;` before working on a new problem



```
Console
-----
AMPL
ampl: commands grid_lp_batch;
CPLEX 12.6.1.0: optimal solution; objective 2000
4 dual simplex iterations (1 in phase I)
x_1 = 0.2

x_2 = 0.4

x_3 = 0.4

T = 2000

ampl:
```

Motivating example 2: Cloud computing

- Now, each resource further charges a set up cost of fixed amount
 - Resource 1: Set up cost = 200 dollars
 - Resource 2: Set up cost = 100 dollars
 - Resource 3: Set up cost = 50 dollars
- Again, the computation job has the following requirements:
 - Requires 10^7 million cycles
 - Completion time $\leq 4,800$ sec
 - Cost $\leq 1,500$ dollars
 - Minimize the completion time
- Problem faced by the job:
 - From which resource should we buy the computing power?
 - How many cycles to buy from each chosen resource?

Yes-or-no decision

- What is the cost of buying cycles from a chosen resource?
- Yes-or-no questions: Is Resource i chosen?
 - E.g. is Resource 3 chosen?
 - **Yes** $\Rightarrow x_3 > 0 \Rightarrow \text{Cost} = 2000 \times x_3 + 200 \times 1$
 - **No** $\Rightarrow x_3 = 0 \Rightarrow \text{Cost} = 2000 \times x_3 + 200 \times 0$
- This type of optimization problems involves 0-or-1 logical decisions among other decisions
- You will learn in the rest of this lecture
 - How to formulate this type of optimization problems
 - How to solve this type of optimization problems

Formulating optimization problem

■ Given:

- n resources
- Resource i offers computing power at a speed of p_i million cycles/sec with cost c_i dollars/sec
- Set up cost for using Resource i is s_i dollars
- Customer requires N million cycles
- Completion time $\leq T_{\max}$
- Cost $\leq C_{\max}$

■ Decision variables:

$$y_i = \begin{cases} 1 & \text{if Resource } i \text{ is chosen} \\ 0 & \text{otherwise} \end{cases}$$

$$x_i = \text{fraction of the job allocated to Resource } i$$

$$T = \text{completion time, which is } \max_i \frac{Nx_i}{p_i}$$

- Exercise: Express the cost of using Resource i in terms of c_i , N , x_i , p_i , s_i , y_i



Formulating optimization problem (cont.)

- Cost:

$$C = \sum_{i=1}^n \left(\frac{c_i N x_i}{p_i} + s_i y_i \right)$$

- Constraint: Cannot have cycles from Resource i if it is not chosen

$$x_i \leq y_i, \quad i = 1, 2, \dots, n$$

- Note that we require $x_i \geq 0$, thus

- Exercise: If $y_i = 0$, what value(s) can x_i take?

- Exercise: If $y_i = 1$, what value(s) can x_i take?

Formulating optimization problem (cont.)

- The problem formulation is

$$\min T$$

subject to

$$T \geq \frac{Nx_i}{p_i}, \quad i = 1, 2, \dots, n$$

$$T \leq T_{\max}$$

$$\sum_{i=1}^n \left(\frac{c_i Nx_i}{p_i} + s_i y_i \right) \leq C_{\max}$$

$$x_i \leq y_i, \quad i = 1, 2, \dots, n$$

$$\sum_{i=1}^n x_i = 1$$

$$x_i \geq 0, \quad i = 1, 2, \dots, n$$

$$y_i \in \{0, 1\}, \quad i = 1, 2, \dots, n$$

Integer programming

- LP in which some or all decision variables can only take non-negative integer values
- Specific type of integer programming (IP):
 - Mixed integer programming (MIP): Some decision variables are integers, others are real
 - Pure integer programming: All decision variables are integers
 - Binary integer programming: All decision variables must be either 0 or 1

LP relaxation

- Relaxing the integer constraints may not give you the right solution



IP problem (IP1)

$$\begin{array}{ll}\min & x_1 + x_2 \\ \text{subject to} & 2x_1 + x_2 \geq 4 \\ & x_1 + 4x_2 \geq 6 \\ & x_1, x_2 \geq 0, \text{ integer}\end{array}$$

LP relaxation of (IP1)

$$\begin{array}{ll}\min & x_1 + x_2 \\ \text{subject to} & 2x_1 + x_2 \geq 4 \\ & x_1 + 4x_2 \geq 6 \\ & x_1, x_2 \geq 0\end{array}$$

Feasible solutions



- Question: Can we round the LP relaxation solution to obtain a solution to (IP1)?

Results and observations

- Rounding the LP relaxation solution may produce an infeasible solution for the original IP problem
- However, for minimization problems, it is always true that:
Optimal value of an IP problem \geq Optimal value of its LP relaxation
- If the LP relaxation gives an integer solution
 - The optimal solution to the LP relaxation is also an optimal solution to the original problem
 - This is true for some special classes of IP problems e.g. network flow problems

Solving IP exactly

- Complete enumeration will require a lot of computation
- A smarter way is to use the branch-and-bound method (you are encouraged to read Winston sections 9.3, 9.4 in your own time)
- In principle, the branch-and-bound method can find the optimal solution but practically it may take a very long time
 - Some IP problems with 60 variables may take hours to run

AMPL/CPLEX for solving example 2

- This is saved in the file `grid_mip.mod`

```
set COMP;      #Set of resources
param c {COMP}; #Cost
param p {COMP}; #Speed
param s {COMP}; #Set up cost
param Tmax;
param Cmax;
param N;
var x {i in COMP} >= 0;
var y {i in COMP} binary;
var T >= 0;
minimize time: T;
subject to T_constraint {i in COMP}: T >= N * x[i] / p[i];
subject to Tmax_constraint: T <= Tmax;
subject to Cmax_constraint: sum {i in COMP}
    (N * c[i] * x[i] / p[i] + s[i] * y[i]) <= Cmax;
subject to x_sum: sum {i in COMP} x[i] = 1;
subject to restriction {i in COMP}: x[i] <= y[i];
```

AMPL/CPLEX for solving example 2 (cont.)

■ Results from the solver

CPLEX 12.6.0.0: optimal integer solution; objective 3333.33333

5 MIP simplex iterations

0 branch-and-bound nodes

x [*] :=

COMP_1 0.333333

COMP_2 0.666667

COMP_3 0

;

y [*] :=

COMP_1 1

COMP_2 1

COMP_3 0

;

T = 3333.33

sum{i in COMP} (N*c[i]*x[i]/p[i] + s[i]*y[i]) = 1466.67

Acknowledgment

- Grid computing example based on Menascé and Casalicchio, “QoS in computing”, IEEE Internet Computing, pp. 85–87, Jul./Aug. 2004.